COMP SUPERSCALAR

COMPSs at BSC

Supercomputers Manual

VERSION: 2.3.RC1808

August 1, 2018

Barcelona
Supercomputing
Center
Centro Nacional de Supercomputación

This manual only provides information about the COMPSs usage at MareNostrum. Specifically, it details the available COMPSs modules, how to load them and how to create and track COMPSs jobs.

If you want to install COMPSs on your local machine please refer to the *COMPSs Installation Manual* available at our webpage `http://compss.bsc.es`.

For further information about the application's execution please refer to the *COMPSs User Manual: Application execution guide* available at `http://compss.bsc.es` .

For further information about the application's development please refer to the *COMPSs User Manual: Application development guide* available at `http://compss.bsc.es/` .

For full COMPSs example application (codes, execution commands, results, logs, etc.) please refer to the *COMPSs Sample Applications* available at `http://compss.bsc.es/` .

Contents

List of Figures

iii

1 COMP Superscalar (COMPSs)

COMP Superscalar (COMPSs) is a programming model which aims to ease the development of applications for distributed infrastructures, such as Clusters, Grids and Clouds. COMP Superscalar also features a runtime system that exploits the inherent parallelism of applications at execution time.

For the sake of programming productivity, the COMPSs model has four key characteristics:

- **Sequential programming:** COMPSs programmers do not need to deal with the typical duties of parallelization and distribution, such as thread creation and synchronization, data distribution, messaging or fault tolerance. Instead, the model is based on sequential programming, which makes it appealing to users that either lack parallel programming expertise or are looking for better programmability.

- **Infrastructure unaware:** COMPSs offers a model that abstracts the application from the underlying distributed infrastructure. Hence, COMPSs programs do not include any detail that could tie them to a particular platform, like deployment or resource management. This makes applications portable between infrastructures with diverse characteristics.

- **Standard programming languages:** COMPSs is based on the popular programming language Java, but also offers language bindings for Python and C/C++ applications. This facilitates the learning of the model, since programmers can reuse most of their previous knowledge.

- **No APIs:** In the case of COMPSs applications in Java, the model does not require to use any special API call, pragma or construct in the application; everything is pure standard Java syntax and libraries. With regard the Python and C/C++ bindings, a small set of API calls should be used on the COMPSs applications.

2 Common usage

2.1 Available COMPSs modules

COMPSs is configured as a Linux Module. As shown in next Figure, the users can type the `module available COMPSs` command to list the supported COMPSs modules in the supercomputer. The users can also execute the `module load COMPSs/<version>` command to load an specific COMPSs module.

```
$ module available COMPSs
---------- /apps/modules/modulefiles/tools ----------
COMPSs/1.3
COMPSs/1.4
COMPSs/2.0
COMPSs/2.1
COMPSs/2.2

COMPSs/release(default)
COMPSs/trunk


$ module load COMPSs/release
load java/1.8.0u66 (PATH, MANPATH, JAVA_HOME, JAVA_ROOT, JAVA_BINDIR,
                    SDK_HOME, JDK_HOME, JRE_HOME)
load MKL/11.0.1 (LD_LIBRARY_PATH)
load PYTHON/2.7.3 (PATH, MANPATH, LD_LIBRARY_PATH, C_INCLUDE_PATH)
load COMPSs/release (PATH, MANPATH, COMPSS_HOME)
```

The following command can be run to check if the correct COMPSs version has been loaded:

```
$ enqueue_compss --version
COMPSs version <version>
```

2.2 Configuration

The COMPSs module contains **all** the COMPSs dependencies, including Java, Python and MKL. Modifying any of these dependencies can cause execution failures and thus, we **do not** recomend to change them. Before running any COMPSs job please check your environment and, if needed, comment out any line inside the `.bashrc` file that loads custom COMPSs, Java, Python and/or MKL modules.

The COMPSs module needs to be loaded in all the nodes that will run COMPSs jobs. Consequently, the `module load` **must** be included in your `.bashrc` file. To do so, please run the following command with the corresponding COMPSs version:

```
$ cat "module load COMPSs/release" >> ~/.bashrc
```

Log out and back in again to check that the file has been correctly edited. The next listing shows an example of the output generated by well loaded COMPSs installation.

```
$ exit
$ ssh USER@SC
load java/1.8.0u66 (PATH, MANPATH, JAVA_HOME, JAVA_ROOT, JAVA_BINDIR,
                    SDK_HOME, JDK_HOME, JRE_HOME)
load MKL/11.0.1 (LD_LIBRARY_PATH)
load PYTHON/2.7.3 (PATH, MANPATH, LD_LIBRARY_PATH, C_INCLUDE_PATH)
load COMPSs/release (PATH, MANPATH, COMPSS_HOME)

$ enqueue_compss --version
COMPSs version <version>
```

> Please remember that COMPSs runs in several nodes and your current
> enviroment is not exported to them. Thus, all the needed environment
> variables **must** be loaded through the *.bashrc* file.

> Please remember that PyCOMPSs uses Python 2.7 by default. In
> order to use Python 3, the Python 2.7 module **must** be unloaded after
> loading COMPSs module, and then load the Python 3 module.

2.3 COMPSs Job submission

COMPSs jobs can be easily submited by running the **enqueue_compss** command. This
command allows to configure any **runcompss** option and some particular queue options
such as the queue system, the number of nodes, the wallclock time, the master working
directory, the workers working directory and number of tasks per node.

Next, we provide detailed information about the *enqueue_compss* command:

```
$ enqueue_compss -h

Usage: enqueue_compss [queue_system_options] [COMPSs_options]
          application_name [application_arguments]

* Options:

  General:

    --help, -h                          Print this help message

  Queue system configuration:

    --sc_cfg=<name>                     SuperComputer configuration file to use.
                                        Must exist inside queues/cfgs/
                                        Default: default

  Submission configuration:

    --exec_time=<minutes>               Expected execution time of the application (in minutes)
                                        Default: 10
```

```
    --num_nodes=<int>                     Number of nodes to use
                                          Default: 2

    --num_switches=<int>                  Maximum number of different switches.
                                          Select 0 for no restrictions.
                                          Maximum nodes per switch: 18
                                          Only available for at least 4 nodes.
                                          Default: 0

   --queue=<name>                         Queue name to submit the job. Depends on the queue system.
                                          For example (Nord3): bsc_cs | bsc_debug | debug | interactive
                                          Default: default

   --reservation=<name>                   Reservation to use when submitting the job.
                                          Default: disabled

   --constraints=<constraints>            Constraints to pass to queue system.
                                          Default: disabled
   --qos=<qos>                            Quality of Service to pass to the queue system.
                                          Default: default

   --job_dependency=<jobID>               Postpone job execution until the job dependency has ended.
                                          Default: None

   --storage_home=<string>                Root installation dir of the storage implementation
                                          Default: null

   --storage_props=<string>               Absolute path of the storage properties file
                                          Mandatory if storage_home is defined

Launch configuration:

  --cpus_per_node=<int>                   Available CPU computing units on each node
                                          Default: 48

  --gpus_per_node=<int>                   Available GPU computing units on each node
                                          Default: 0

  --max_tasks_per_node=<int>              Maximum number of simultaneous tasks running on a node
                                          Default: -1

  --node_memory=<MB>                      Maximum node memory: disabled | <int> (MB)
                                          Default: disabled

  --network=<name>                        Communication network for transfers:
                                          default | ethernet | infiniband | data.
                                          Default: infiniband

  --prolog="<string>"                     Task to execute before launching COMPSs (Notice the quotes)
                                          If the task has arguments split them by "," rather than spaces.
                                          This argument can appear multiple times for more than one
                                          prolog action
                                          Default: Empty

  --epilog="<string>"                     Task to execute after executing the COMPSs application (Notice
                                          the quotes)
                                          If the task has arguments split them by "," rather than spaces.
                                          This argument can appear multiple times for more than one
                                          epilog action
                                          Default: Empty

  --master_working_dir=<path>             Working directory of the application
                                          Default: .

  --worker_working_dir=<name | path>      Worker directory. Use: scratch | gpfs | <path>
                                          Default: scratch

  --worker_in_master_cpus=<int>           Maximum number of CPU computing units that the master node can
                                          run as worker. Cannot exceed cpus_per_node.
```

```
                                     Default: 24

  --worker_in_master_memory=<int> MB      Maximum memory in master node assigned to the worker. Cannot
                                          exceed the node_memory.
                                          Mandatory if worker_in_master_cpus is specified.
                                          Default: 50000

  --jvm_worker_in_master_opts="<string>"  Extra options for the JVM of the COMPSs Worker in the Master
                                          Node.
                                          Each option separed by "," and without blank spaces (Notice the
                                          quotes)
                                          Default:

  --container_image=<path>                Runs the application by means of a container engine image
                                          Default: Empty

  --container_compss_path=<path>          Path where compss is installed in the container image
                                          Default: /opt/COMPSs

  --container_opts="<string>"             Options to pass to the container engine
                                          Default: empty

  --elasticity=<max_extra_nodes>          Activate elasticity specifiying the maximum extra nodes (ONLY
                                          AVAILABLE FORM SLURM CLUSTERS WITH NIO ADAPTOR)
                                          Default: 0

Runcompss configuration:

Tools enablers:

  --graph=<bool>, --graph, -g             Generation of the complete graph (true/false)
                                          When no value is provided it is set to true
                                          Default: false

  --tracing=<level>, --tracing, -t        Set generation of traces and/or tracing level
                                          ( [ true | basic ] | advanced | false)
                                          True and basic levels will produce the same traces.
                                          When no value is provided it is set to true
                                          Default: false

  --monitoring=<int>, --monitoring, -m    Period between monitoring samples (milliseconds)
                                          When no value is provided it is set to 2000
                                          Default: 0

  --external_debugger=<int>,
  --external_debugger                     Enables external debugger connection on the specified port
                                          (or 9999 if empty)
                                          Default: false

Runtime configuration options:

  --task_execution=<compss|storage>       Task execution under COMPSs or Storage.
                                          Default: compss

  --storage_conf=<path>                   Path to the storage configuration file
                                          Default: None

  --project=<path>                        Path to the project XML file
                                          Default: /apps/COMPSs/2.3/Runtime/configuration/xml/projects/
                                          default_project.xml

  --resources=<path>                      Path to the resources XML file
                                          Default: /apps/COMPSs/2.3/Runtime/configuration/xml/resources/
                                          default_resources.xml

  --lang=<name>                           Language of the application (java/c/python)
                                          Default: Inferred is possible. Otherwise: java

  --summary                               Displays a task execution summary at the end of the application
                                          execution
```

```
                                          Default: false

  --log_level=<level>, --debug, -d        Set the debug level: off | info | debug
                                          Default: off

Advanced options:

  --extrae_config_file=<path>             Sets a custom extrae config file. Must be in a shared disk
                                          between all COMPSs workers.
                                          Default: null

  --comm=<ClassName>                      Class that implements the adaptor for communications
                                          Supported adaptors: es.bsc.compss.nio.master.NIOAdaptor
                                                            | es.bsc.compss.gat.master.GATAdaptor
                                          Default: es.bsc.compss.nio.master.NIOAdaptor

  --conn=<className>                      Class that implements the runtime connector for the cloud
                                          Supported connectors:
                                                      es.bsc.compss.connectors.DefaultSSHConnector
                                                    | es.bsc.compss.connectors.DefaultNoSSHConnector
                                          Default: es.bsc.compss.connectors.DefaultSSHConnector

  --scheduler=<className>                 Class that implements the Scheduler for COMPSs
                                          Supported schedulers:
                                            es.bsc.compss.scheduler.fullGraphScheduler.FullGraphScheduler
                                          | es.bsc.compss.scheduler.fifoScheduler.FIFOScheduler
                                          | es.bsc.compss.scheduler.resourceEmptyScheduler.
                                            ResourceEmptyScheduler
                                          Default: es.bsc.compss.scheduler.loadBalancingScheduler.
                                                  LoadBalancingScheduler

  --scheduler_config_file=<path>          Path to the file which contains the scheduler configuration.
                                          Default: Empty

  --library_path=<path>                   Non-standard directories to search for libraries (e.g. Java JVM
                                          library, Python library, C binding library)
                                          Default: Working Directory

  --classpath=<path>                      Path for the application classes / modules
                                          Default: Working Directory

  --appdir=<path>                         Path for the application class folder.
                                          Default: /home/user/

  --pythonpath=<path>                     Additional folders or paths to add to the PYTHONPATH
                                          Default: /home/user/

  --base_log_dir=<path>                   Base directory to store COMPSs log files (a .COMPSs/ folder
                                          will be created inside this location)
                                          Default: User home

  --specific_log_dir=<path>               Use a specific directory to store COMPSs log files (the folder
                                          MUST exist and no sandbox is created)
                                          Warning: Overwrites --base_log_dir option
                                          Default: Disabled

  --uuid=<int>                            Preset an application UUID
                                          Default: Automatic random generation

  --master_name=<string>                  Hostname of the node to run the COMPSs master
                                          Default:

  --master_port=<int>                     Port to run the COMPSs master communications.
                                          Only for NIO adaptor
                                          Default: [43000,44000]

  --jvm_master_opts="<string>"            Extra options for the COMPSs Master JVM. Each option separed
                                          by "," and without blank spaces (Notice the quotes)
                                          Default:
```

```
    --jvm_workers_opts="<string>"        Extra options for the COMPSs Workers JVMs. Each option separed
                                         by "," and without blank spaces (Notice the quotes)
                                         Default: -Xms1024m,-Xmx1024m,-Xmn400m

    --cpu_affinity="<string>"            Sets the CPU affinity for the workers
                                         Supported options: disabled, automatic, user defined map of
                                         the form "0-8/9,10,11/12-14,15,16"
                                         Default: automatic

    --gpu_affinity="<string>"            Sets the GPU affinity for the workers
                                         Supported options: disabled, automatic, user defined map of
                                         the form "0-8/9,10,11/12-14,15,16"
                                         Default: automatic

    --task_count=<int>                   Only for C/Python Bindings. Maximum number of different
                                         functions/methods, invoked from the application, that have
                                         been selected as tasks
                                         Default: 50

    --input_profile=<path>               Path to the file which stores the input application profile
                                         Default: Empty

    --output_profile=<path>              Path to the file to store the application profile at the end of
                                         the execution
                                         Default: Empty

    --PyObject_serialize=<bool>          Only for Python Binding. Enable the object serialization to
                                         string when possible (true/false).
                                         Default: false

    --persistent_worker_c=<bool>         Only for C Binding. Enable the persistent worker in c
                                         (true/false).
                                         Default: false

    --enable_external_adaptation=<bool>  Enable external adaptation. This option will disable the
                                         Resource Optimizer.
                                         Default: false

* Application name:

    For Java applications:   Fully qualified name of the application
    For C applications:      Path to the master binary
    For Python applications: Path to the .py file containing the main program

* Application arguments:

    Command line arguments to pass to the application. Can be empty.
```

3 MareNostrum 4

3.1 Basic queue commands

The MareNostrum supercomputer uses the SLURM (Simple Linux Utility for Resource Management) workload manager. The basic commands to manage jobs are listed below:

- **sbatch** Submit a batch job to the SLURM system

- **scancel** Kill a running job

- **squeue -u \<username\>** See the status of jobs in the SLURM queue

For more extended information please check the *SLURM: Quick start user guide* at `https://slurm.schedmd.com/quickstart.html` .

3.2 Tracking COMPSs jobs

When submitting a COMPSs job a temporal file will be created storing the job information. For example:

```
$ enqueue_compss \
  --exec_time=15 \
  --num_nodes=3 \
  --cpus_per_node=16 \
  --master_working_dir=. \
  --worker_working_dir=gpfs \
  --lang=python \
  --log_level=debug \
  <APP> <APP_PARAMETERS>


SC Configuration:          default.cfg
Queue:                     default
Reservation:               disabled
Num Nodes:                 3
Num Switches:              0
GPUs per node:             0
Job dependency:            None
Exec-Time:                 00:15
Storage Home:              null
Storage Properties:        null
Other:
        --sc_cfg=default.cfg
        --cpus_per_node=48
        --master_working_dir=.
        --worker_working_dir=gpfs
        --lang=python
        --classpath=.
        --library_path=.
        --comm=es.bsc.compss.nio.master.NIOAdaptor
        --tracing=false
        --graph=false
        --pythonpath=.
        <APP> <APP_PARAMETERS>
Temp submit script is: /scratch/tmp/tmp.pBG5yfFxEo

$ cat /scratch/tmp/tmp.pBG5yfFxEo
#!/bin/bash
#
#SBATCH --job-name=COMPSs
```

```
#SBATCH --workdir=.
#SBATCH -o compss-%J.out
#SBATCH -e compss-%J.err
#SBATCH -N 3
#SBATCH -n 144
#SBATCH --exclusive
#SBATCH -t00:15:00
...
```

In order to trac the jobs state users can run the following command:

```
$ squeue
JOBID     PARTITION  NAME    USER  TIME_LEFT  TIME_LIMIT  START_TIME  ST NODES  CPUS  NODELIST
474130     main     COMPSs   XX    0:15:00     0:15:00       N/A       PD   3    144   -
```

The specific COMPSs logs are stored under the `~/.COMPSs/` folder; saved as a local *runcompss* execution. For further details please check *COMPSs User Manual: Application Execution* available at our webpage `http://compss.bsc.es` .

4   MinoTauro

4.1   Basic queue commands

The MinoTauro supercomputer uses the SLURM (Simple Linux Utility for Resource Management) workload manager. The basic commands to manage jobs are listed below:

- **sbatch** Submit a batch job to the SLURM system

- **scancel** Kill a running job

- **squeue -u <username>** See the status of jobs in the SLURM queue

For more extended information please check the *SLURM: Quick start user guide* at `https://slurm.schedmd.com/quickstart.html` .

4.2   Tracking COMPSs jobs

When submitting a COMPSs job a temporal file will be created storing the job information. For example:

```
$ enqueue_compss \
  --exec_time=15 \
  --num_nodes=3 \
  --cpus_per_node=16 \
  --master_working_dir=. \
  --worker_working_dir=gpfs \
  --lang=python \
  --log_level=debug \
  <APP> <APP_PARAMETERS>


SC Configuration:       default.cfg
Queue:                  default
Reservation:            disabled
Num Nodes:              3
Num Switches:           0
GPUs per node:          0
Job dependency:         None
Exec-Time:              00:15
Storage Home:           null
Storage Properties:     null
Other:
        --sc_cfg=default.cfg
        --cpus_per_node=16
        --master_working_dir=.
        --worker_working_dir=gpfs
        --lang=python
        --classpath=.
        --library_path=.
        --comm=es.bsc.compss.nio.master.NIOAdaptor
        --tracing=false
        --graph=false
        --pythonpath=.
        <APP> <APP_PARAMETERS>
Temp submit script is: /scratch/tmp/tmp.pBG5yfFxEo

$ cat /scratch/tmp/tmp.pBG5yfFxEo
#!/bin/bash
#
#SBATCH --job-name=COMPSs
```

```
#SBATCH --workdir=.
#SBATCH -o compss-%J.out
#SBATCH -e compss-%J.err
#SBATCH -N 3
#SBATCH -n 48
#SBATCH --exclusive
#SBATCH -t00:15:00
...
```

In order to trac the jobs state users can run the following command:

```
$ squeue
JOBID  PARTITION   NAME    USER  ST  TIME    NODES  NODELIST (REASON)
XXXX   projects    COMPSs  XX    R   00:02       3  nvb[6-8]
```

The specific COMPSs logs are stored under the `~/.COMPSs/` folder; saved as a local *runcompss* execution. For further details please check *COMPSs User Manual: Application Execution* available at our webpage `http://compss.bsc.es` .

5 Nord 3

5.1 Basic queue commands

The Nord3 supercomputer uses the LSF (Load Sharing Facility) workload manager. The basic commands to manage jobs are listed below:

- **bsub** Submit a batch job to the LSF system

- **bkill** Kill a running job

- **bjobs** See the status of jobs in the LSF queue

- **bqueues** Information about LSF batch queues

For more extended information please check the *IBM Platform LSF Command Reference* at `https://www.ibm.com/support/knowledgecenter/en/SSETD4_9.1.2/lsf_kc_cmd_ref.html` .

5.2 Tracking COMPSs jobs

When submitting a COMPSs job a temporal file will be created storing the job information. For example:

```
$ enqueue_compss \
  --exec_time=15 \
  --num_nodes=3 \
  --cpus_per_node=16 \
  --master_working_dir=. \
  --worker_working_dir=gpfs \
  --lang=python \
  --log_level=debug \
  <APP> <APP_PARAMETERS>


SC Configuration:        default.cfg
Queue:                   default
Reservation:             disabled
Num Nodes:               3
Num Switches:            0
GPUs per node:           0
Job dependency:          None
Exec-Time:               00:15
Storage Home:            null
Storage Properties:      null
Other:
        --sc_cfg=default.cfg
        --cpus_per_node=16
        --master_working_dir=.
        --worker_working_dir=gpfs
        --lang=python
        --classpath=.
        --library_path=.
        --comm=es.bsc.compss.nio.master.NIOAdaptor
        --tracing=false
        --graph=false
        --pythonpath=.
        <APP> <APP_PARAMETERS>
Temp submit script is: /scratch/tmp/tmp.pBG5yfFxEo
```

```
$ cat /scratch/tmp/tmp.pBG5yfFxEo
#!/bin/bash
#
#BSUB -J COMPSs
#BSUB -cwd .
#BSUB -oo compss-%J.out
#BSUB -eo compss-%J.err
#BSUB -n 3
#BSUB -R "span[ptile=1]"
#BSUB -W 00:15
...
```

In order to trac the jobs state users can run the following command:

```
$ bjobs
JOBID   USER    STAT   QUEUE   FROM_HOST   EXEC_HOST   JOB_NAME   SUBMIT_TIME
XXXX    bscXX   PEND   XX      login1      XX          COMPSs     Month Day Hour
```

The specific COMPSs logs are stored under the `~/.COMPSs/` folder; saved as a local *runcompss* execution. For further details please check *COMPSs User Manual: Application Execution* available at our webpage `http://compss.bsc.es` .

6 Enabling COMPSs Monitor

6.1 Configuration

As supercomputer nodes are connection restricted, the better way to enable the *COMPSs Monitor* is from the users local machine. To do so please install the following packages:

- COMPSs Runtime

- COMPSs Monitor

- sshfs

For further details about the COMPSs packages installation and configuration please refer to the *COMPSs Installation Manual* available at our webpage `http://compss.bsc.es` . If you are not willing to install COMPSs in your local machine please consider to download our Virtual Machine available at our webpage.

Once the packages have been installed and configured, users need to mount the sshfs directory as follows. The `SC_USER` stands for your supercomputer's user, the `SC_ENDPOINT` to the supercomputer's public endpoint and the `TARGET_LOCAL_FOLDER` to the local folder where you wish to deploy the supercomputer files):

```
compss@bsc:~$ scp $HOME/.ssh/id_dsa.pub ${SC_USER}@mn1.bsc.es:~/id_dsa_local.pub
compss@bsc:~$ ssh SC_USER@SC_ENDPOINT
               "cat ~/id_dsa_local.pub >> ~/.ssh/authorized_keys;
               rm ~/id_dsa_local.pub"
compss@bsc:~$ mkdir -p TARGET_LOCAL_FOLDER/.COMPSs
compss@bsc:~$ sshfs -o IdentityFile=$HOME/.ssh/id_dsa -o allow_other
               SC_USER@SC_ENDPOINT:~/.COMPSs
               TARGET_LOCAL_FOLDER/.COMPSs
```

Whenever you wish to unmount the sshfs directory please run:

```
compss@bsc:~$ sudo umount TARGET_LOCAL_FOLDER/.COMPSs
```

6.2 Execution

Access the COMPSs Monitor through its webpage (`http://localhost:8080/compss-monitor` by default) and log in with the `TARGET_LOCAL_FOLDER` to enable the COMPSs Monitor for MareNostrum.

Please remember that to enable **all** the COMPSs Monitor features applications must be ran with the *-m* flag. For further information please check the *COMPSs User Manual: Application Execution* available at our webpage `http://compss.bsc.es` .

Figure 1 illustrates how to login and Figure 2 shows the COMPSs Monitor main page for an application run inside a Supercomputer.

Figure 1: COMPSs Monitor login for Supercomputers

Figure 2: COMPSs Monitor main page for a test application at Supercomputers

Please find more details on the COMPSs framework at
http://compss.bsc.es